



Modern IoT Stack

M1 Computer Science

Docker, MQTT, InfluxDB, Grafana & Node-RED

Complete Architecture for Professional IoT Applications



Joseph Azar

Associate Professor - Université de Franche-Comté

Researcher FEMTO-ST

The Problem: Traditional IoT Setup

Naive Approach (System Installation)

1. Manually install Mosquitto
2. Manually install InfluxDB
3. Manually install Grafana
4. Configure each service individually
5. Manage dependencies

Scenario: Direct installation on the system 6. Resolve port conflicts

-  "It works on my machine" syndrome
-  Different configurations per OS
-  Incompatible dependencies
-  Setup takes hours

Problems: •  Difficult to clean up

✓ Modern Solution: Docker

Scenario: Docker Compose

```
# One single command  
docker-compose up
```

```
# Everything starts:  
# ✓ Mosquitto  
# ✓ InfluxDB  
# ✓ Grafana  
# ✓ Node-RED
```



- ⚡ Setup in 30 seconds
- ⏪ Reproducible everywhere
- 🧹 Easy cleanup
- 🔒 Service isolation
- 📦 Precise versioning

Benefits:

- 🌐 Automatic networking



What is Docker?

Simple Definition: Docker = "Magic boxes" for your applications

Analogy: Shipping Containers

- **Standardized:** Same format everywhere
- **Portable:** Ship, train, truck
- **Isolated:** Protected content
- **Inventoried:** Know what's inside

```
# Image = "recipe" of container
FROM node:18

# What's inside the container
WORKDIR /app
COPY . .
RUN npm install

# How to start it
CMD [ "node", "app.js" ]
```

- **Image:** The "blueprint" of the container
- **Container:** Running instance
- **Isolated:** Own filesystem
- **Lightweight:** Shares OS kernel

Docker vs Virtual Machines

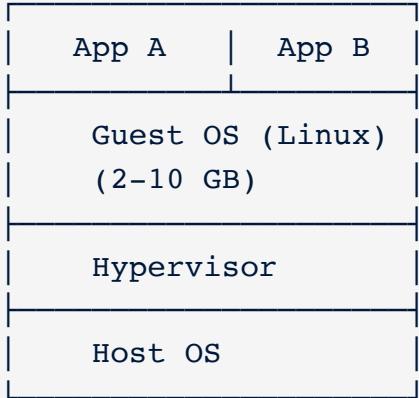
Feature	Virtual Machine (VM)	Docker Container
OS	Full OS (Windows, Linux...)	Uses host OS kernel (no full OS)
Size	 Several GB (2-10 GB)	 Few MB (10-500 MB)
Startup	 Minutes (1-5 min)	 Seconds (< 5s)
Isolation	 Complete (virtual hardware)	 Process-level (namespaces)
Performance	 Significant overhead	 Near native
Use Case	Different OS, max security	Microservices, CI/CD, dev

 **For IoT:** Docker is perfect - lightweight, fast to deploy, and easily orchestrates multiple services

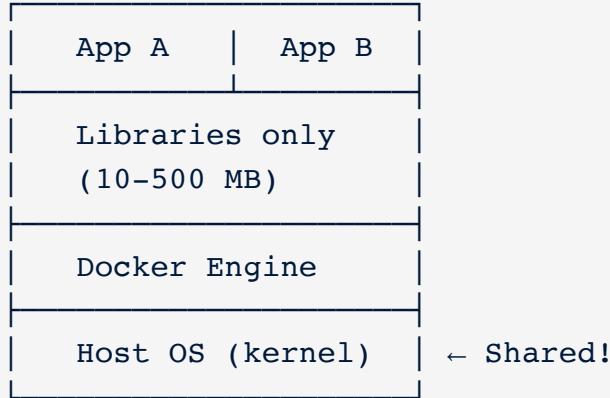
VM vs Docker: Architecture Explained

 **Key Difference:** VM = Entire computer simulated | Docker = Just the application + libraries

Virtual Machine:



Docker Container:



Virtual Machine:

- Includes entire Guest OS
- Duplicates OS for each VM
- Heavy (GBs per VM)
- Slow boot (minutes)

Docker Container:

- Shares host OS kernel
- Only app + dependencies
- Lightweight (MBs per container)
- Fast boot (seconds)



Images vs Containers

Class Analogy: Image = Class, Container = Instance

Docker Image

- 📋 **Template** - read-only
- 🎉 **Layers** - stacked
- 📦 **Immutable** - never changes
- 🏭 **Build** once, run anywhere

```
# List images  
docker images
```

```
# Download an image  
docker pull nginx:latest
```

```
# Build an image  
docker build -t myapp:1.0 .
```

Docker Container

- ⏴ **Instance** - running
- 🖊 **Writable** - can be modified
- ⚡ **Ephemeral** - can be destroyed
- 📈 **State** - running, stopped, paused

```
# List active containers  
docker ps
```

```
# Create and run container  
docker run -d -p 80:80 nginx:latest
```

```
# Stop a container  
docker stop my_container
```



Image vs Container Example



- One **image** can spawn multiple **containers**
- Containers are **isolated** from each other
- Each container has its own **writable layer**

Key Points: • Image remains **unchanged** when containers run



Docker Volumes: Data Persistence

⚠ Problem: When you delete a container, ALL its data is lost!

✗ Without Volume

```
# Launch InfluxDB
docker run influxdb:1.8

# Create data
# ... write to DB

# Delete the container
docker rm -f container_id

# 😱 All data is gone!
```





Docker Volumes: Solution

✓ With Volume

```
# Launch with volume
docker run -v influxdb_data:/var/lib/influxdb influxdb:1.8

# Create data
# ... write to DB

# Delete the container
docker rm -f container_id

# ✓ Data still there!
# Relaunch with same volume
docker run -v influxdb_data:/var/lib/influxdb influxdb:1.8
```



Result: Persistent data survives container restarts!

Types of Docker Volumes

Type	Syntax	Location	Use Case
Named Volume	volumeName:/path	Managed by Docker	 Persistent data (DB)
Bind Mount	./host/path:/container/path	Absolute path on host	 Config, source code
tmpfs	--tmpfs /path	RAM memory	 Temporary data

Concrete Examples

```
# docker-compose.yml
services:
  influxdb:
    volumes:
      - influxdb_data:/var/lib/influxdb # Named volume

  mosquitto:
    volumes:
      - ./mosquitto.conf:/mosquitto/config/mosquitto.conf # Bind mount

volumes:
  influxdb_data: # Volume declaration
```



What Happens When You Delete?

- **Container**: Deleted with its internal data
 - **Volume**: Remains intact and can be reused
- When you delete a container:** • **Image**: Still there, unchanged

```
# Example workflow
docker run -v mydata:/data myapp      # Create container with volume
# ... work with data ...
docker rm -f myapp_container          # Delete container

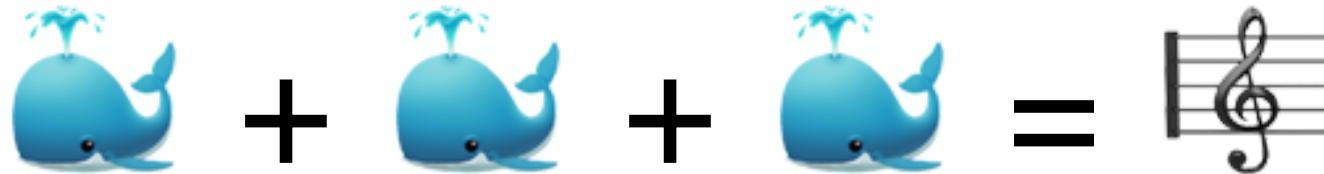
# Volume "mydata" still exists!
docker volume ls                      # See all volumes
docker run -v mydata:/data myapp      # Reattach to same data
```



Best Practice: Always use volumes for important data (databases, configs, logs)

Docker Compose

Orchestrating Multiple Containers Made Easy



What is Docker Compose?

Definition: Tool to define and run multi-container applications with a YAML file

Without Docker Compose

```
docker network create iot-network
docker run -d --name mosquitto --network iot-network -p 1883:1883 \
    -v ./mosquitto.conf:/mosquitto/config/mosquitto.conf eclipse-mosquitto
docker run -d --name influxdb --network iot-network -p 8086:8086 \
    -v influxdb_data:/var/lib/influxdb influxdb:1.8
docker run -d --name grafana --network iot-network -p 3000:3000 \
    -v grafana_data:/var/lib/grafana grafana/grafana
docker run -d --name nodered --network iot-network -p 1880:1880 \
    -v nodered_data:/data nodered/node-red
```



Problems: Complex, repetitive, error-prone

✓ With Docker Compose

```
# docker-compose.yml
services:
mosquitto:
  image: eclipse-mosquitto
  ports:
    - "1883:1883"
  volumes:
    - ./mosquitto.conf:/mosquitto/config/mosquitto.conf

influxdb:
  image: influxdb:1.8
  ports:
    - "8086:8086"
  volumes:
    - influxdb_data:/var/lib/influxdb

grafana:
  image: grafana/grafana
  ports:
    - "3000:3000"
  volumes:
    - grafana_data:/var/lib/grafana

nodered:
  image: nodered/node-red
  ports:
    - "1880:1880"
  volumes:
    - nodered_data:/data

volumes:
  influxdb_data:
```



```
grafana_data:  
nodered_data:
```

```
# One command!  
docker-compose up -d
```



Essential Docker Compose Commands

Command	Description	Use Case
<code>docker-compose up</code>	Create and start all services	 First time or after changes
<code>docker-compose up -d</code>	Start in background (detached)	 Free up terminal
<code>docker-compose down</code>	Stop and remove containers	 Cleanup (keeps volumes)
<code>docker-compose down -v</code>	Stop and remove volumes too	 Complete reset
<code>docker-compose ps</code>	List active containers	 Monitoring
<code>docker-compose logs</code>	View logs of all services	 Debugging
<code>docker-compose build</code>	Rebuild images	 After Dockerfile change
<code>docker-compose restart</code>	Restart services	 After config change

Docker Compose: Automatic Networking

Magic: Docker Compose automatically creates a private network for your services!

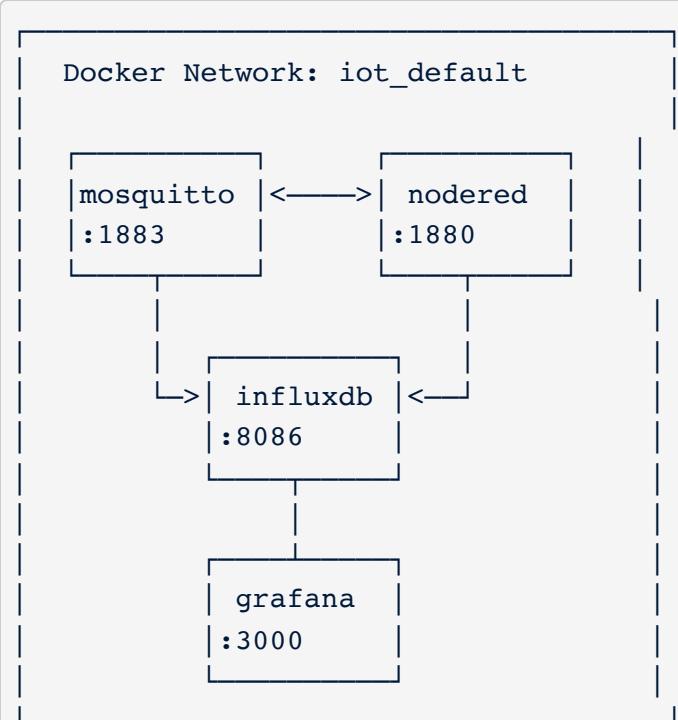
```
services:  
mosquitto:  
  image: eclipse-mosquitto  
  # Accessible via "mosquitto:1883"  
  
nodered:  
  image: nodered/node-red  
  depends_on:  
    - mosquitto  
  # Node-RED can connect via: mqtt://mosquitto:1883
```



-  Service name = hostname
-  Isolated from outside world
-  Fast internal communication

 **Benefits:** •  Built-in DNS

Network Diagram



| Port mapping
▼

Host Machine

localhost:1883 → mosquitto:1883
localhost:3000 → grafana:3000

// In Node.js code running inside Docker Compose
const mqtt = require('mqtt');

// ✅ CORRECT - use service name
const client = mqtt.connect('mqtt://mosquitto:1883');

Docker Cleanup: Prune Commands

⚠ Problem: Docker accumulates data (images, containers, volumes) that takes up space

Command	What it cleans	⚠ Warning
<code>docker container prune</code>	Stopped containers	 Safe, only stopped ones
<code>docker image prune</code>	Unused images (dangling)	 Safe, keeps tagged images
<code>docker image prune -a</code>	ALL unused images	 Removes tagged images too
<code>docker volume prune</code>	Unattached volumes	 DATA LOSS possible!
<code>docker system prune</code>	Containers, networks, images	 Removes many things
<code>docker system prune -a --volumes</code>	EVERYTHING unused	 COMPLETE RESET - Danger!

```
# Check disk space used by Docker  
docker system df
```





MQTT: IoT Protocol

Messaging for the Internet of Things



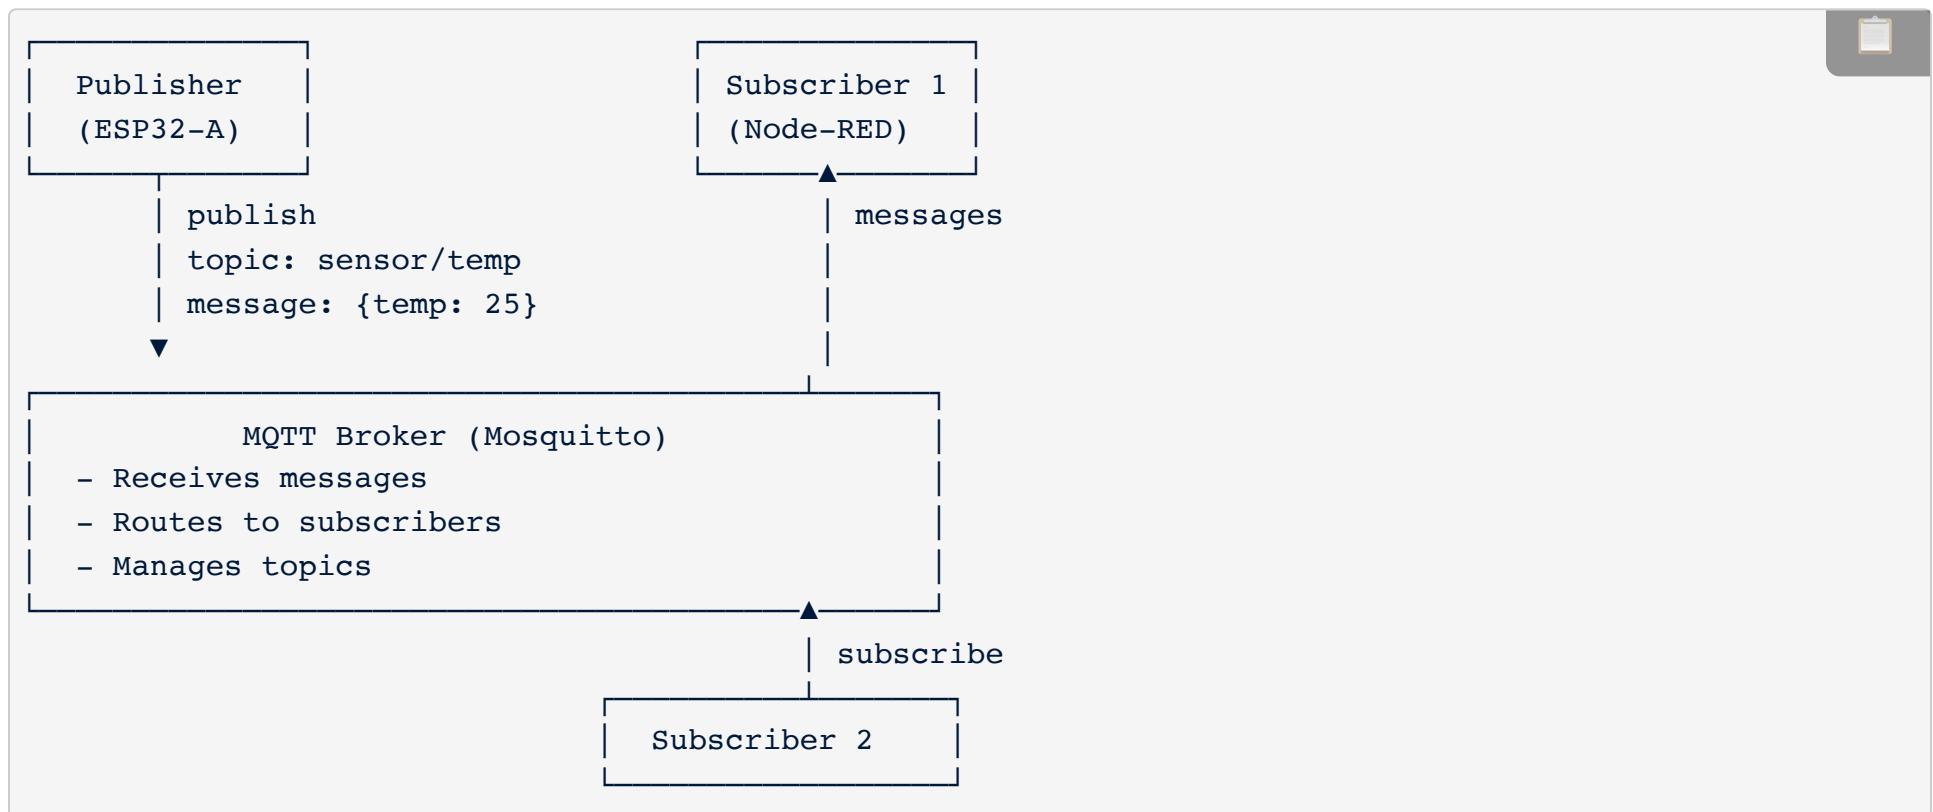


What is MQTT?

MQTT = Message Queuing Telemetry Transport

- 📬 Lightweight **messaging** protocol
- 🌐 Designed for **IoT** and unstable networks
- 🛰 **Publish/Subscribe** pattern
- ⚡ Very low power consumption (battery friendly)

🏛️ MQTT Architecture





MQTT Roles



Three Main Roles:

- **Publisher:** Sends messages to a topic
- **Subscriber:** Receives messages from a topic
- **Broker:** Central server that routes messages



Topic Structure

```
home/livingroom/temperature  
home/livingroom/humidity  
home/bedroom/temperature  
sensor/esp32/a  
sensor/esp32/b  
sensor/esp32/led
```



Wildcards:

```
home/+/temperature      # + = one level  
sensor/#                 # # = all levels
```



Publish/Subscribe Pattern

- **Decoupling**: Publisher and Subscriber don't know each other
- **Topics**: Communication channels
- **Many-to-Many**: 1 publisher → N subscribers

Key Principles: • **Asynchronous**: Non-blocking

Example: Temperature Sensor

```
// Publisher (ESP32 - sensor)
const mqtt = require('mqtt');
const client = mqtt.connect('mqtt://mosquitto:1883');

client.on('connect', () => {
  setInterval(() => {
    const temp = Math.random() * 30 + 20;
    const payload = JSON.stringify({
      temperature: temp.toFixed(2),
      humidity: 65,
      timestamp: Date.now()
    });
    client.publish('sensor/esp32/a', payload);
    console.log('Published:', payload);
  }, 10000); // Every 10s
});
```





MQTT Subscriber Example

```
// Subscriber (Dashboard)
const mqtt = require('mqtt');
const client = mqtt.connect('mqtt://mosquitto:1883');

client.on('connect', () => {
  // Subscribe to topic
  client.subscribe('sensor/esp32/a', (err) => {
    if (!err) {
      console.log('Subscribed to sensor/esp32/a');
    }
  });
});

client.on('message', (topic, message) => {
  const data = JSON.parse(message.toString());
  console.log('Received:', data);
  // Display on dashboard, save to DB...
});
```



- ESP32 doesn't need to know who's listening
- Multiple subscribers can receive the same message



Benefits: • New subscribers can join anytime

MQTT vs HTTP: Fundamental Differences

Feature	HTTP (Request/Response)	MQTT (Pub/Sub)
Pattern	Client-Server (1:1)	Publish-Subscribe (M:N)
Connection	Short, per request	Persistent (WebSocket)
Overhead	 Large headers (~200 bytes)	 Very light (~2 bytes)
Real-time	 Polling needed	 Instant push
Bandwidth	 High	 Minimal
Battery	 Drains quickly	 Very efficient
Unstable Network	 Frequent timeouts	 QoS, auto-reconnect

🎯 Scenario 1: Temperature Monitoring - HTTP

✗ HTTP Approach (Polling)

```
// Dashboard must poll every 5 seconds
setInterval(async () => {
  const response = await fetch('http://esp32/temp');
  const data = await response.json();
  updateUI(data);
}, 5000);
```



- **17,280 requests / day**
 - 12 requests/minute × 60 min × 24h
 - Constant network traffic
- **Latency (up to 5s before update)**
 - Temperature changes at 10:00:01
 - Dashboard updates at 10:00:05 (next poll)
 - Not real-time!
- **ESP32 battery drained**
 - HTTP server always listening
 - Processing 17K+ requests daily
 - No sleep mode possible
- **Bandwidth wasted**
 - Most polls return same data
 - HTTP headers overhead (~200 bytes each)
 - ~3.5 MB/day just for headers

Problems with HTTP Polling:

🎯 Scenario 1: Temperature Monitoring - MQTT

✓ MQTT Approach (Push)

```
// Publisher (ESP32) - sends only when temperature changes
client.on('connect', () => {
    const currentTemp = readSensor();

    // Only publish if change > 0.5°C
    if (Math.abs(currentTemp - lastTemp) > 0.5) {
        client.publish('sensor/temp',
            JSON.stringify({temp: currentTemp}));
        lastTemp = currentTemp;
    }
});

// Subscriber (Dashboard) - receives instantly
client.on('message', (topic, message) => {
    const data = JSON.parse(message);
    updateUI(data); // Updates in < 100ms
});
```



Benefits with MQTT Push:

- ⚡ **Few messages/day:** ~10-50 messages (vs 17,280 HTTP) = 99.7% reduction
- ⚡ **Real-time:** Updates in < 100ms (vs up to 5s latency with polling)
- 🔋 **Battery preserved:** ESP32 sleeps between readings, months vs days
- 📡 **Optimal bandwidth:** ~100 bytes/day (vs 3.5 MB) = 35,000x efficient

🎯 Scenario 2: LED Control - HTTP Approach

✗ HTTP Approach

```
// ESP32 must be an HTTP server
app.post('/led/on', (req, res) => {
    digitalWrite(LED_PIN, HIGH);
    res.json({status: 'on'});
});
```



- **ESP32 needs fixed IP address**
 - Dashboard must know exact IP
 - Difficult with DHCP networks
- **Firewall / NAT complications**
 - ESP32 behind router not accessible
 - Need port forwarding configuration
- **No auto-reconnect**
 - Lost connection = manual restart
 - No built-in recovery mechanism
- **Heavy server code on ESP32**
 - HTTP server library = large memory footprint
 - Complex request parsing

Problems with HTTP:

🎯 Scenario 2: LED Control - MQTT Approach

✓ MQTT Approach

```
// ESP32 is just a subscriber
const mqtt = require('mqtt');
const client = mqtt.connect('mqtt://mosquitto:1883');

client.subscribe('sensor/esp32/led');

client.on('message', (topic, message) => {
  const state = message.toString() === 'true';
  digitalWrite(LED_PIN, state ? HIGH : LOW);
  console.log(`LED turned ${state ? 'ON' : 'OFF'}`);
});

// Dashboard publishes (from anywhere)
client.publish('sensor/esp32/led', 'true');
```



-  **No server on ESP32**
 - ESP32 is a client (subscriber)
 - Minimal memory footprint
-  **Works behind NAT**
 - ESP32 connects OUT to broker
 - No firewall/port issues
-  **Auto-reconnect**
 - Built-in reconnection logic
 - Resilient to network issues
-  **Minimal code**
 - Simple subscribe + callback
 - Lightweight MQTT library

Benefits with MQTT:

🤔 Why MQTT for IoT?

✓ Key Advantages for IoT

1. Power Efficiency

- Persistent connection (no repeated handshake)
- Ultra-light messages (2 bytes header vs 200+ HTTP)
- Intelligent keep-alive
- → **Batteries last months/years**

2. Unstable Networks

- QoS (Quality of Service): 0, 1, 2
- Messages persisted if client disconnected
- Automatic reconnection
- → **Works even with unstable 3G/4G**

MQTT Scalability

- 1 broker can handle millions of clients
- Hierarchical topics with wildcards
- No direct coupling publisher-subscriber

3. Scalability • Add sensors without modifying code

Concrete Numbers

Metric	HTTP	MQTT
Minimum message size	~200 bytes	~2 bytes
Connection setup	Every request	Once at startup
Notification latency	Polling dependent (5-60s)	< 100ms (real-time)
Battery consumption	 High	 Minimal
Bandwidth (1000 msg/h)	~200 KB	~2 KB



Mosquitto: The MQTT Broker

Mosquitto = Open-source MQTT broker, lightweight and performant (Eclipse Foundation)



Basic Configuration

```
# mosquitto.conf
listener 1883
allow_anonymous true

# For production:
# listener 1883
# allow_anonymous false
# password_file /mosquitto/config/passwd
```



In Docker Compose

```
services:
mosquitto:
  image: eclipse-mosquitto
  container_name: mosquitto
  ports:
    - "1883:1883" # MQTT
  volumes:
    - ./mosquitto.conf:/mosquitto/config/mosquitto.conf
```



🧪 Testing Mosquitto

```
# Start the stack
docker-compose up -d

# Subscriber in one terminal
docker exec -it mosquitto \
mosquitto_sub -h localhost -t "test/#"

# Publisher in another terminal
docker exec -it mosquitto \
mosquitto_pub -h localhost \
-t "test/hello" \
-m "Hello MQTT!"

# The subscriber receives:
# Hello MQTT!
```



📊 Monitoring

```
# View all messages (debug)
docker exec -it mosquitto \
mosquitto_sub -h localhost -t "#" -v

# System topics (stats)
docker exec -it mosquitto \
mosquitto_sub -h localhost -t "\$SYS/#"
```





Visual Programming for IoT



What is Node-RED?

Node-RED = Flow-based visual programming tool for IoT

- **Drag & Drop**: No coding needed (or almost!)
- **Nodes**: Reusable blocks (MQTT, HTTP, DB...)
- **Web Interface**: Browser accessible
- **Node.js**: Built on Node.js (extensible)

🎯 Use Cases

- **MQTT Routing**: sensor → database
- **Transformation**: JSON → InfluxDB format
- **Logic**: If temp > 30, send alert
- **Dashboard**: Real-time visualization
- **Integration**: APIs, webhooks, services

:green_diamond: Node Types

Input	MQTT in, HTTP in, Inject
Output	MQTT out, HTTP, Debug
Function	Function, Switch, Change
Storage	InfluxDB, MySQL, File
Dashboard	Chart, Gauge, Text, Button

Node-RED in Docker Compose

```
services:  
  nodered:  
    image: nodered/node-red  
    container_name: nodered  
    ports:  
      - "1880:1880"  
    volumes:  
      - nodered_data:/data # Persist flows  
    depends_on:  
      - mosquitto  
    environment:  
      - TZ=Europe/Paris  
  
volumes:  
  nodered_data:
```



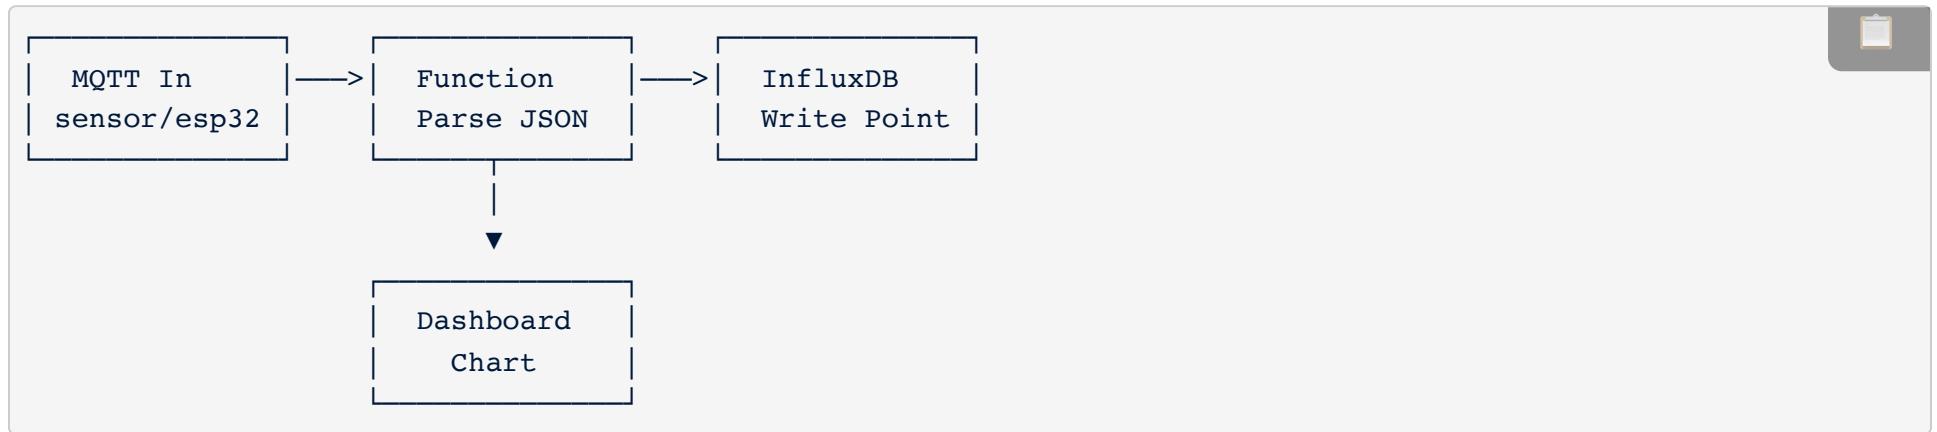
 **Interface:** `http://localhost:1880` after `docker-compose up`



Node-RED Flow Example

Goal: Receive MQTT data → Save to InfluxDB → Display on dashboard

Visual Flow



Function Node Code

```
// Parse received JSON
const data = JSON.parse(msg.payload);

// Format for InfluxDB
msg.payload = {
  measurement: 'temperature',
  fields: {
    value: parseFloat(data.temperature),
    humidity: parseFloat(data.humidity)
  },
  tags: { sensor: 'esp32_a' },
  timestamp: new Date()
```

```
};  
return msg;
```



Time Series Database





What is InfluxDB?

InfluxDB = Database specialized for **time series** (temporal data)

🤔 What is a Time Series?

Time Series = Data with timestamp



Timestamp	Value
2024-01-15 10:00:00	22.5
2024-01-15 10:01:00	22.7
2024-01-15 10:02:00	22.6
2024-01-15 10:03:00	23.1

Examples:

- 📈 Sensor temperature
- 💰 Stock prices
- 🌐 Network traffic
- 🏁 Server metrics

🎯 Why InfluxDB and not SQL/MongoDB?

SQL (MySQL, Postgres)	<ul style="list-style-type: none">✗ Not optimized for time series✗ Slow on temporal aggregations✗ Inefficient storage✗ No native downsampling
NoSQL (MongoDB)	<ul style="list-style-type: none">✗ No temporal compression✗ Verbose temporal queries✗ No retention policies⚠ Good but not specialized
InfluxDB	<ul style="list-style-type: none">✓ 100x compression (time series)✓ Ultra-fast aggregations✓ Auto retention policies✓ Continuous queries✓ InfluxQL language (SQL-like)



InfluxDB Concepts

Database: iot_data

```
|  
|   └─ Measurement: temperature  
|       └─ Tags (indexed)  
|           └─ sensor = "esp32_a"  
|           └─ location = "livingroom"  
|  
|       └─ Fields (values)  
|           └─ value = 22.5  
|           └─ humidity = 65  
|  
└─ Timestamp: 2024-01-15T10:00:00Z
```



SQL Analogy:

- Database = Database
- Measurement = Table
- Tags = Indexed columns
- Fields = Columns
- Timestamp = Primary key



Writing Data to InfluxDB

Line Protocol Format

```
measurement,tag1=value1,tag2=value2 field1=value1,field2=value2 timestamp  
# Example  
temperature,sensor=esp32_a,room=living value=22.5,humidity=65 161071200000000000
```



Node.js Example

```
const Influx = require('influx');
const client = new Influx.InfluxDB({
  host: 'influxdb',
  database: 'iot_data'
});

client.writePoints([
  {
    measurement: 'temperature',
    tags: {
      sensor: 'esp32_a',
      location: 'livingroom'
    },
    fields: {
      value: 22.5,
      humidity: 65
    },
    timestamp: new Date()
  }]
);
```



InfluxQL Queries

```
-- Last 24 hours
SELECT * FROM temperature
WHERE time > now() - 24h

-- Average per hour
SELECT MEAN(value) FROM temperature
WHERE time > now() - 7d
GROUP BY time(1h), sensor

-- With tags
SELECT * FROM temperature
WHERE sensor = 'esp32_a'
AND time > now() - 1h
```



Why InfluxDB for IoT?

IoT-Specific Advantages

1. **Temporal Compression**

- Specialized algorithms (Gorilla, etc.)
- 100x compression vs classic SQL
- Example: 1 year of data/minute = 5 MB instead of 500 MB

2. **Aggregation Performance**

- Ultra-fast mean/min/max/percentile calculations
- Continuous Queries: automatic pre-calculation
- Example: "Average temp per hour over 1 year" in < 100ms



Retention Policies (Auto-cleanup)

- Define data lifetime
- Automatic downsampling (1min → 1h → 1d)
- Example: Raw data 7d, aggregated 1 year

```
-- Create retention policy
CREATE RETENTION POLICY "one_week" ON "iot_data"
  DURATION 7d
  REPLICATION 1
  DEFAULT;

-- Continuous Query for downsampling
CREATE CONTINUOUS QUERY "cq_30m" ON "iot_data"
BEGIN
  SELECT MEAN(value) AS mean_value
  INTO "iot_data"."one_year"."temperature_30m"
  FROM "temperature"
  GROUP BY time(30m), sensor
END;

-- Result:
-- Raw data: 7 days (auto-deleted)
-- 30min averages: 1 year (for charts)
```



InfluxDB with Docker

```
services:  
  influxdb:  
    image: influxdb:1.8  
    container_name: influxdb-server  
    ports:  
      - "8086:8086"  
    volumes:  
      - influxdb_data:/var/lib/influxdb  
    environment:  
      - INFLUXDB_DB=iot_data  
      - INFLUXDB_ADMIN_USER=admin  
      - INFLUXDB_ADMIN_PASSWORD=supersecret  
  
volumes:  
  influxdb_data:
```

Test via CLI

```
# Enter container  
docker exec -it influxdb-server influx  
  
# Inside InfluxDB CLI  
> USE iot_data  
> SELECT * FROM temperature ORDER BY time DESC LIMIT 5
```



Visualization and Monitoring





What is Grafana?

Grafana = Open-source visualization and monitoring platform

- **Dashboards:** Charts, gauges, tables...
- **Data Sources:** InfluxDB, Prometheus, MySQL, etc.
- **Alerting:** Automatic notifications
- **Customizable:** Themes, plugins, panels

🎯 Use Cases

- **IoT Dashboards:** Real-time sensors
- **Infrastructure:** CPU, RAM, network
- **Business Metrics:** KPIs, analytics
- **Industrial:** Production, quality

📊 Visualization Types

- Time Series (line, area)
- Gauge, progress bar
- Stat (single value)
- Table
- Heatmap
- Pie Chart, donut
- Geomap

Grafana with Docker



```
services:  
  grafana:  
    image: grafana/grafana  
    container_name: grafana  
    ports:  
      - "3000:3000"  
    volumes:  
      - grafana_data:/var/lib/grafana  
    environment:  
      - GF_SECURITY_ADMIN_PASSWORD=admin  
  
volumes:  
  grafana_data:
```

Interface:

<http://localhost:3000>

Default: admin / admin

Creating a Grafana Dashboard

Steps: Data Source → Dashboard → Panel → Query → Visualization

1 Configure Data Source

Configuration → Data Sources → Add data source
Type: InfluxDB
URL: http://influxdb:8086
Database: iot_data
Min time interval: 10s



2 Create Panel with Query

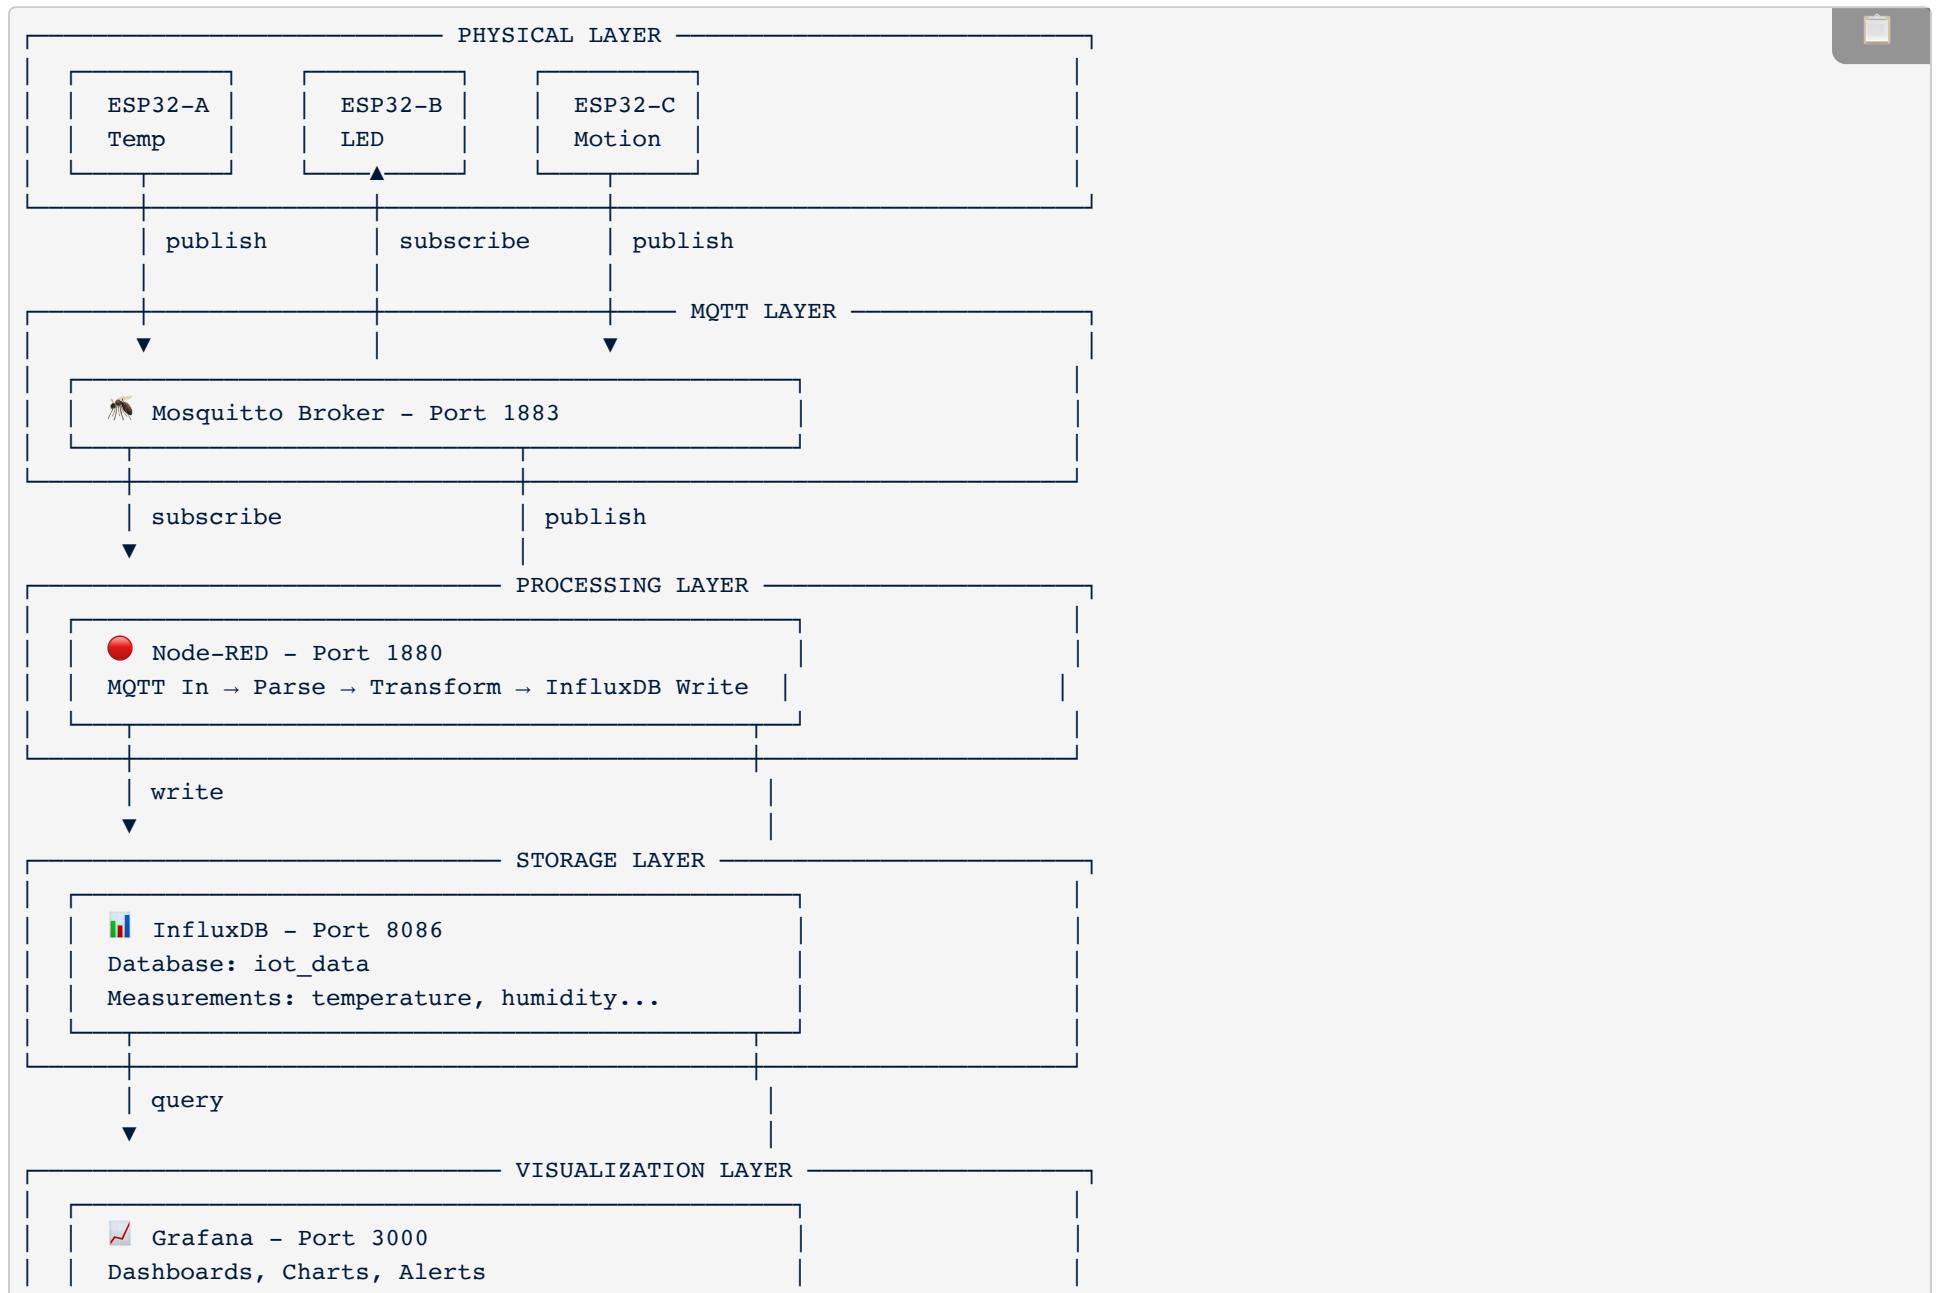
```
-- Real-time temperature
SELECT mean("value")
FROM "temperature"
WHERE $timeFilter
  AND "sensor" = 'esp32_a'
GROUP BY time($__interval) fill(null)

-- Grafana variables:
-- $timeFilter : automatic time filter
-- $__interval : adaptive interval
```





Complete IoT Stack Architecture





Our docker-compose.yml Explained



```
services:
# 🐝 Mosquitto: MQTT Broker
mosquitto:
  image: eclipse-mosquitto           # Official image
  container_name: mosquitto          # Fixed container name
  ports:
    - "1883:1883"                  # MQTT port (host:container)
  volumes:
    - ./mosquitto.conf:/mosquitto/config/mosquitto.conf # Config (bind mount)

# 🔴 Node-RED: Flow processing
nodered:
  image: nodered/node-red
  container_name: nodered
  ports:
    - "1880:1880"                  # Web interface
  volumes:
    - nodered_data:/data          # Flows persisted (named volume)
  depends_on:
    - mosquitto                   # Wait for mosquitto to be up

# 📈 InfluxDB: Time series database
influxdb:
  image: influxdb:1.8               # Version 1.8 (InfluxQL)
  container_name: influxdb-server
  ports:
    - "8086:8086"                 # HTTP API
  volumes:
    - influxdb_data:/var/lib/influxdb # Data persisted

# 📈 Grafana: Visualization
grafana:
  image: grafana/grafana
```

```
container_name: grafana
ports:
  - "3000:3000"                      # Web interface
volumes:
  - grafana_data:/var/lib/grafana     # Dashboards persisted
```

```
# 📁 Named volumes (managed by Docker)
```

```
volumes:
  nodered_data:
  influxdb_data:
  grafana_data:
```

Key Points

 **Automatic Network:** All services can communicate via their names

 **Named Volumes:** Data survives `docker-compose down`

 **Bind Mount:** `mosquitto.conf` mounted from host

 **depends_on:** Startup order (Node-RED after Mosquitto)

🧪 Hands-On: Launch Complete Stack

1 Start the stack

```
cd /path/to/week_1  
docker-compose up -d  
  
# Verify everything is running  
docker-compose ps
```



2 Test MQTT with simulated ESP32

```
# Terminal 1 - Launch publisher (ESP32-A)  
node esp32_a.js  
  
# Terminal 2 - Launch subscriber (ESP32-B)  
node esp32_b_subscriber.js  
  
# Terminal 3 - Monitor all MQTT messages  
docker exec -it mosquitto mosquitto_sub -h localhost -t "#" -v
```



 Access Web Interfaces

Service	URL	Credentials
Node-RED	http://localhost:1880	None (default)
Grafana	http://localhost:3000	admin / admin
InfluxDB API	http://localhost:8086	-

Cleanup and Management

Basic commands

```
# Stop services  
docker-compose stop  
  
# Restart services  
docker-compose start  
  
# Restart specific service  
docker-compose restart mosquitto  
  
# View logs  
docker-compose logs  
docker-compose logs -f mosquitto
```



Cleanup

```
# Stop and remove containers  
docker-compose down  
  
# Remove containers + volumes  
# ⚠️ LOSES ALL DATA  
docker-compose down -v  
  
# Rebuild images  
docker-compose build  
docker-compose up -d --build  
  
# Check disk space  
docker system df
```



⚠️ Attention:

- `docker-compose down`: Removes containers, KEEPS volumes
- `docker-compose down -v`: Removes EVERYTHING (containers + volumes)



Summary

Docker

- Application containerization
- Images vs Containers
- Volumes for persistence
- Docker Compose multi-service
- Automatic networking
- Essential commands

MQTT

- Pub/Sub pattern
- Mosquitto broker
- Topics and wildcards
- MQTT vs HTTP
- Why for IoT

Node-RED

- Visual programming
- Flows and nodes
- MQTT/InfluxDB integration
- Real-time dashboard

InfluxDB & Grafana

- Time Series Database
- Why vs SQL/Mongo
- InfluxQL queries
- Visualization
- Dashboards and alerts



You now master a professional IoT stack!

Docker + MQTT + Node-RED + InfluxDB + Grafana =

Resources and Documentation

Official Documentation

-  **Docker:** docs.docker.com
-  **MQTT:** mqtt.org
-  **Mosquitto:** mosquitto.org
-  **Node-RED:** nodered.org
-  **InfluxDB:** docs.influxdata.com
-  **Grafana:** grafana.com/docs

Training Projects

1.  **Smart Home:** Multi-sensors + dashboard
2.  **Plant Monitor:** Soil, light, temp
3.  **Energy Monitor:** Real-time consumption
4.  **GPS Tracker:** Location + history
5.  **Industrial:** Production monitoring

? Questions?

Feel free to ask your questions!



Joseph Azar

Associate Professor - Université de Franche-Comté
Researcher FEMTO-ST

 joseph.azar@univ-fcomte.fr

